

```
In [2]: import pandas as pd
import numpy as np
```

```
In [3]: from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.seasonal import seasonal_decompose
from pmdarima import auto_arima
```

/Users/mirek/opt/anaconda3/lib/python3.8/site-packages/statsmodels/compat/pandas.py:61: FutureWarning: pandas.Int64Index is deprecated and will be removed from pandas in a future version. Use pandas.Index with the appropriate dtype instead.

```
from pandas import Int64Index as NumericIndex
```

```
In [4]: from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error
```

```
In [5]: df=pd.read_csv('airline_passengers.csv',index_col='Month',parse_dates=True)
df
```

Out[5]: **Thousands of Passengers**

Month	
1949-01-01	112
1949-02-01	118
1949-03-01	132
1949-04-01	129
1949-05-01	121
...	...
1960-08-01	606
1960-09-01	508
1960-10-01	461
1960-11-01	390
1960-12-01	432


144 rows × 1 columns

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 144 entries, 1949-01-01 to 1960-12-01
Data columns (total 1 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Thousands of Passengers              144 non-null   int64
dtypes: int64(1)
memory usage: 2.2 KB
```


```
In [7]: df.index.freq='MS'
```

```
In [8]: df.plot();
```



```
In [9]: seasonalDecomposition=seasonal_decompose(df,model='mul')
```

```
In [10]: seasonalDecomposition.plot();
```



```
In [13]: auto_arima(df,seasonal=True,m=12,trace=True).summary()
```

Performing stepwise search to minimize aic

```
ARIMA(2,1,2)(1,1,1)[12] : AIC=1020.048, Time=0.48 sec
ARIMA(0,1,0)(0,1,0)[12] : AIC=1031.508, Time=0.01 sec
ARIMA(1,1,0)(1,1,0)[12] : AIC=1020.393, Time=0.05 sec
ARIMA(0,1,1)(0,1,1)[12] : AIC=1021.003, Time=0.07 sec
ARIMA(2,1,2)(0,1,1)[12] : AIC=1019.935, Time=0.28 sec
ARIMA(2,1,2)(0,1,0)[12] : AIC=1019.290, Time=0.11 sec
ARIMA(2,1,2)(1,1,0)[12] : AIC=1019.546, Time=0.28 sec
ARIMA(1,1,2)(0,1,0)[12] : AIC=1024.160, Time=0.06 sec
ARIMA(2,1,1)(0,1,0)[12] : AIC=1017.847, Time=0.11 sec
ARIMA(2,1,1)(1,1,0)[12] : AIC=1017.914, Time=0.25 sec
ARIMA(2,1,1)(0,1,1)[12] : AIC=1018.359, Time=0.24 sec
ARIMA(2,1,1)(1,1,1)[12] : AIC=1018.248, Time=0.47 sec
ARIMA(1,1,1)(0,1,0)[12] : AIC=1022.393, Time=0.03 sec
ARIMA(2,1,0)(0,1,0)[12] : AIC=1022.393, Time=0.03 sec
ARIMA(3,1,1)(0,1,0)[12] : AIC=1019.084, Time=0.11 sec
ARIMA(1,1,0)(0,1,0)[12] : AIC=1020.393, Time=0.02 sec
ARIMA(3,1,0)(0,1,0)[12] : AIC=1023.666, Time=0.05 sec
ARIMA(3,1,2)(0,1,0)[12] : AIC=1021.083, Time=0.20 sec
ARIMA(2,1,1)(0,1,0)[12] intercept : AIC=inf, Time=0.26 sec
```

Best model: ARIMA(2,1,1)(0,1,0)[12]

Total fit time: 3.138 seconds

Out [13]:

## SARIMAX Results

Dep. Variable:	y	No. Observations:	144			
Model:	SARIMAX(2, 1, 1)x(0, 1, [], 12)	Log Likelihood	-504.923			
Date:	Tue, 26 Apr 2022	AIC	1017.847			
Time:	06:45:11	BIC	1029.348			
Sample:	0	HQIC	1022.520			
	- 144					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.5960	0.085	6.987	0.000	0.429	0.763
ar.L2	0.2143	0.091	2.343	0.019	0.035	0.394
ma.L1	-0.9819	0.038	-25.602	0.000	-1.057	-0.907
sigma2	129.3150	14.557	8.883	0.000	100.784	157.846
Ljung-Box (L1) (Q):	0.00	Jarque-Bera (JB):	7.68			
Prob(Q):	0.98	Prob(JB):	0.02			
Heteroskedasticity (H):	2.33	Skew:	-0.01			
Prob(H) (two-sided):	0.01	Kurtosis:	4.19			

## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

In [14]:

```
train=df.iloc[:-12]
test=df.iloc[-12:]
len(test)
```

Out [14]: 12

In [15]:

```
model=SARIMAX(train,order=(2,1,1),seasonal_order=(0,1,0,12))
results=model.fit()
```

## RUNNING THE L-BFGS-B CODE

\* \* \*

Machine precision = 2.220D-16

N = 4 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 3.42129D+00 |proj g|= 2.27842D-01

At iterate 5 f= 3.39217D+00 |proj g|= 1.21741D-02

At iterate 10 f= 3.39006D+00 |proj g|= 2.06863D-03

At iterate 15 f= 3.39003D+00 |proj g|= 1.15951D-03

\* \* \*

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

\* \* \*

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
4	18	32	1	0	0	1.039D-05	3.390D+00

F = 3.3900246510399326

CONVERGENCE: REL\_REDUCTION\_OF\_F\_&lt;=\_FACTR\*EPSMCH

/Users/mirek/opt/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/stat  
 espace/sarimax.py:966: UserWarning: Non-stationary starting autoregressive  
 parameters found. Using zeros as starting parameters.

warn('Non-stationary starting autoregressive parameters')

/Users/mirek/opt/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/stat  
 espace/sarimax.py:978: UserWarning: Non-invertible starting MA parameters f  
 ound. Using zeros as starting parameters.

warn('Non-invertible starting MA parameters found.')

This problem is unconstrained.

Warning: more than 10 function and gradient  
 evaluations in the last line search. Termination  
 may possibly be caused by a bad search direction.

In [16]: `results.summary()`

Out [16]:

## SARIMAX Results

**Dep. Variable:** Thousands of Passengers **No. Observations:** 132

**Model:** SARIMAX(2, 1, 1)x(0, 1, [], 12) **Log Likelihood** -447.483

**Date:** Tue, 26 Apr 2022 **AIC** 902.967

**Time:** 06:45:41 **BIC** 914.083

**Sample:** 01-01-1949 **HQIC** 907.481  
- 12-01-1959

**Covariance Type:** opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.5253	0.929	-0.566	0.572	-2.345	1.295
ar.L2	0.0110	0.275	0.040	0.968	-0.529	0.551
ma.L1	0.3064	0.935	0.328	0.743	-1.527	2.140
sigma2	108.0093	13.498	8.002	0.000	81.553	134.466

**Ljung-Box (L1) (Q):** 0.00 **Jarque-Bera (JB):** 1.39

**Prob(Q):** 0.97 **Prob(JB):** 0.50

**Heteroskedasticity (H):** 1.47 **Skew:** -0.09

**Prob(H) (two-sided):** 0.23 **Kurtosis:** 3.50

## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


In [17]:

```
start=len(train)
end=start+len(test)-1
predictions=results.predict(start=start,end=end,dynamic=False).rename('SAR')
```

In [18]:

```
ax=test.plot(legend=True,figsize=(12,6))
predictions.plot(legend=True)
```

Out[18]: <AxesSubplot:xlabel='Month'>



In [19]: `mean_absolute_percentage_error(test,predictions)`

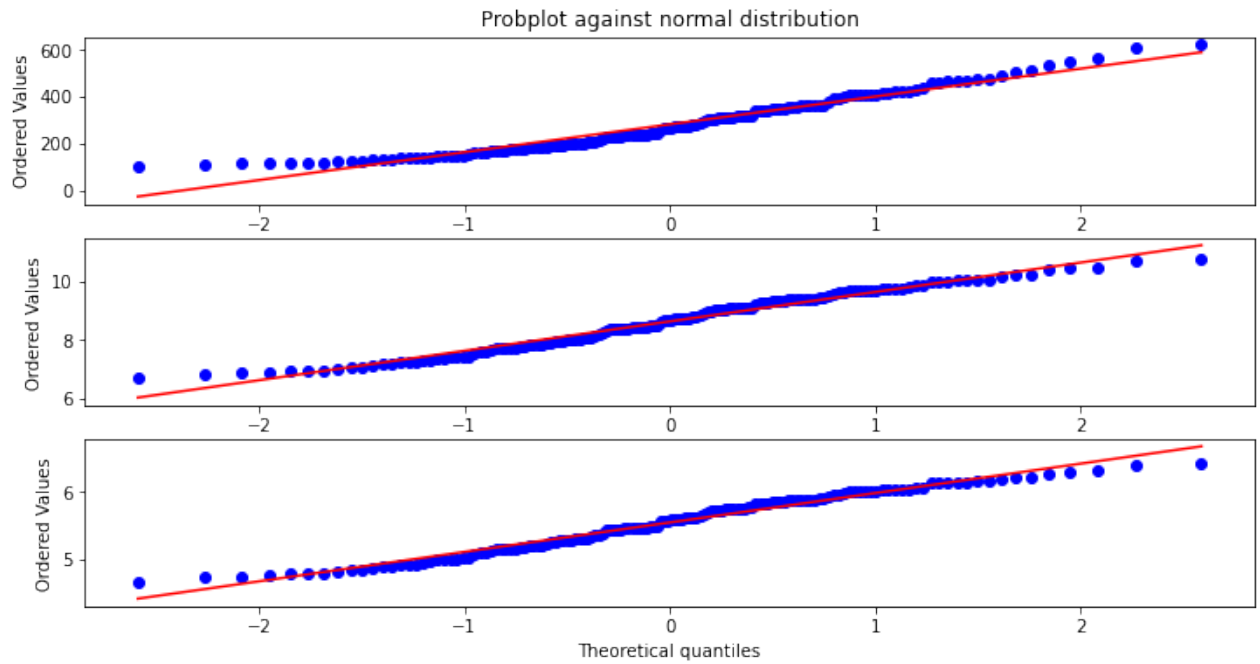
Out[19]: 0.044317548784737405

In [20]: `import scipy.stats as stats`  
`import matplotlib.pyplot as plt`

In [21]: `x=df['Thousands of Passengers']`  
`xt, lambda_mle=stats.boxcox(x)`  
`x_log=np.log(x)`

`fig, (ax1,ax2,ax3)=plt.subplots(3,figsize=(12,6))`  
`stats.probplot(x,dist=stats.norm,plot=ax1);`  
`stats.probplot(xt,dist=stats.norm,plot=ax2);`  
`stats.probplot(x_log,dist=stats.norm,plot=ax3);`  
`ax1.set_title('Probplot against normal distribution')`  
`ax2.set_title('')`  
`ax3.set_title('')`  
`ax1.set_xlabel('')`

Out[21]: Text(0.5, 0, '')



In [22]: lambda\_mle

Out[22]: 0.14802265137037945

In [23]: `df['Log']=np.log(df["Thousands of Passengers"])`  
df

Out[23]:

	Thousands of Passengers	Log
Month		
1949-01-01	112	4.718499
1949-02-01	118	4.770685
1949-03-01	132	4.882802
1949-04-01	129	4.859812
1949-05-01	121	4.795791
...	...	...
1960-08-01	606	6.406880
1960-09-01	508	6.230481
1960-10-01	461	6.133398
1960-11-01	390	5.966147
1960-12-01	432	6.068426

144 rows x 2 columns



In [24]:

```
auto_arma(df['Log'],seasonal=True,m=12,trace=True).summary()
```

Performing stepwise search to minimize aic

```
ARIMA(2,0,2)(1,1,1)[12] intercept : AIC=-480.365, Time=0.81 sec
ARIMA(0,0,0)(0,1,0)[12] intercept : AIC=-358.005, Time=0.03 sec
ARIMA(1,0,0)(1,1,0)[12] intercept : AIC=-472.836, Time=0.28 sec
ARIMA(0,0,1)(0,1,1)[12] intercept : AIC=-423.080, Time=0.30 sec
ARIMA(0,0,0)(0,1,0)[12] : AIC=-152.756, Time=0.01 sec
ARIMA(2,0,2)(0,1,1)[12] intercept : AIC=-485.489, Time=0.64 sec
ARIMA(2,0,2)(0,1,0)[12] intercept : AIC=-455.362, Time=0.29 sec
ARIMA(2,0,2)(0,1,2)[12] intercept : AIC=-483.632, Time=1.63 sec
ARIMA(2,0,2)(1,1,0)[12] intercept : AIC=-479.075, Time=0.74 sec
ARIMA(2,0,2)(1,1,2)[12] intercept : AIC=-482.918, Time=1.32 sec
ARIMA(1,0,2)(0,1,1)[12] intercept : AIC=-486.812, Time=0.54 sec
ARIMA(1,0,2)(0,1,0)[12] intercept : AIC=-455.568, Time=0.29 sec
ARIMA(1,0,2)(1,1,1)[12] intercept : AIC=-477.167, Time=0.45 sec
ARIMA(1,0,2)(0,1,2)[12] intercept : AIC=-484.812, Time=1.31 sec
ARIMA(1,0,2)(1,1,0)[12] intercept : AIC=-480.056, Time=0.58 sec
ARIMA(1,0,2)(1,1,2)[12] intercept : AIC=-483.074, Time=1.67 sec
ARIMA(0,0,2)(0,1,1)[12] intercept : AIC=-455.593, Time=0.41 sec
ARIMA(1,0,1)(0,1,1)[12] intercept : AIC=-489.188, Time=0.50 sec
ARIMA(1,0,1)(0,1,0)[12] intercept : AIC=-457.177, Time=0.12 sec
ARIMA(1,0,1)(1,1,1)[12] intercept : AIC=-487.318, Time=0.61 sec
ARIMA(1,0,1)(0,1,2)[12] intercept : AIC=-487.300, Time=1.13 sec
ARIMA(1,0,1)(1,1,0)[12] intercept : AIC=-482.595, Time=0.47 sec
ARIMA(1,0,1)(1,1,2)[12] intercept : AIC=-485.198, Time=1.18 sec
ARIMA(1,0,0)(0,1,1)[12] intercept : AIC=-482.026, Time=0.27 sec
ARIMA(2,0,1)(0,1,1)[12] intercept : AIC=-487.314, Time=0.55 sec
ARIMA(0,0,0)(0,1,1)[12] intercept : AIC=-367.061, Time=0.10 sec
ARIMA(2,0,0)(0,1,1)[12] intercept : AIC=-489.295, Time=0.37 sec
ARIMA(2,0,0)(0,1,0)[12] intercept : AIC=-458.262, Time=0.11 sec
ARIMA(2,0,0)(1,1,1)[12] intercept : AIC=-487.388, Time=0.66 sec
ARIMA(2,0,0)(0,1,2)[12] intercept : AIC=-487.321, Time=0.68 sec
ARIMA(2,0,0)(1,1,0)[12] intercept : AIC=-482.539, Time=0.90 sec
ARIMA(2,0,0)(1,1,2)[12] intercept : AIC=-486.907, Time=1.21 sec
ARIMA(3,0,0)(0,1,1)[12] intercept : AIC=-487.303, Time=0.52 sec
ARIMA(3,0,1)(0,1,1)[12] intercept : AIC=-485.297, Time=0.61 sec
ARIMA(2,0,0)(0,1,1)[12] : AIC=-481.419, Time=0.33 sec
```

Best model: ARIMA(2,0,0)(0,1,1)[12] intercept

Total fit time: 21.662 seconds

Out [24]:

## SARIMAX Results

Dep. Variable:	y	No. Observations:	144			
Model:	SARIMAX(2, 0, 0)x(0, 1, [1], 12)	Log Likelihood	249.648			
Date:	Tue, 26 Apr 2022	AIC	-489.295			
Time:	06:46:57	BIC	-474.881			
Sample:	0	HQIC	-483.438			
	- 144					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
intercept	0.0194	0.008	2.340	0.019	0.003	0.036
ar.L1	0.5750	0.075	7.632	0.000	0.427	0.723
ar.L2	0.2616	0.084	3.122	0.002	0.097	0.426
ma.S.L12	-0.5549	0.106	-5.225	0.000	-0.763	-0.347
sigma2	0.0013	0.000	8.856	0.000	0.001	0.002
Ljung-Box (L1) (Q):	0.02	Jarque-Bera (JB):	6.47			
Prob(Q):	0.90	Prob(JB):	0.04			
Heteroskedasticity (H):	0.58	Skew:	0.07			
Prob(H) (two-sided):	0.07	Kurtosis:	4.08			

## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

In [25]:

```
train=df.iloc[:-12]
test=df.iloc[-12:]
len(test)
```

Out [25]: 12

In [27]:

```
model2=SARIMAX(train['Log'],order=(2,0,0),seasonal_order=(0,1,1,12))
results2=model2.fit()
results2.summary()
```

RUNNING THE L-BFGS-B CODE

\* \* \*

Machine precision = 2.220D-16

N = 4 M = 10

```

At X0          0 variables are exactly at the bounds

At iterate    0      f= -1.60346D+00      |proj g|=  1.42556D+00
At iterate    5      f= -1.67277D+00      |proj g|=  4.61090D+00
At iterate   10      f= -1.68643D+00      |proj g|=  7.10030D-02
At iterate   15      f= -1.68739D+00      |proj g|=  5.02994D-02
At iterate   20      f= -1.68740D+00      |proj g|=  1.44260D-01
At iterate   25      f= -1.68879D+00      |proj g|=  4.28816D-01
  This problem is unconstrained.
At iterate   30      f= -1.69873D+00      |proj g|=  2.38197D-01
At iterate   35      f= -1.69904D+00      |proj g|=  1.53912D-02

```

\* \* \*

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

\* \* \*

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      4       38       79      2      0      0    1.568D-03  -1.699D+00
F = -1.6990374795045518

```

CONVERGENCE: REL\_REDUCTION\_OF\_F\_<=\_FACTR\*EPSMCH

```

Bad direction in the line search;
  refresh the lbfgs memory and restart the iteration.

```

```

Warning: more than 10 function and gradient
  evaluations in the last line search. Termination
  may possibly be caused by a bad search direction.

```

Out [27]:

## SARIMAX Results

<b>Dep. Variable:</b>	Log	<b>No. Observations:</b>	132
<b>Model:</b>	SARIMAX(2, 0, 0)x(0, 1, [1], 12)	<b>Log Likelihood</b>	224.273
<b>Date:</b>	Tue, 26 Apr 2022	<b>AIC</b>	-440.546
<b>Time:</b>	06:47:26	<b>BIC</b>	-429.396
<b>Sample:</b>	01-01-1949	<b>HQIC</b>	-436.018
	- 12-01-1959		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ar.L1</b>	0.6787	0.074	9.183	0.000	0.534	0.824
<b>ar.L2</b>	0.3143	0.076	4.161	0.000	0.166	0.462
<b>ma.S.L12</b>	-0.5626	0.094	-5.992	0.000	-0.747	-0.379
<b>sigma2</b>	0.0013	0.000	8.170	0.000	0.001	0.002

<b>Ljung-Box (L1) (Q):</b>	0.00	<b>Jarque-Bera (JB):</b>	0.98
<b>Prob(Q):</b>	0.97	<b>Prob(JB):</b>	0.61
<b>Heteroskedasticity (H):</b>	0.35	<b>Skew:</b>	0.02
<b>Prob(H) (two-sided):</b>	0.00	<b>Kurtosis:</b>	3.44

## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


In [28]:

```
predictions2=results2.predict(start=start,end=end,dynamic=False).rename('Log')
```

In [29]:

```
ax=test['Log'].plot(legend=True,figsize=(12,6))
predictions2.plot(legend=True)
```

Out[29]: <AxesSubplot:xlabel='Month'>



In [30]: `mean_absolute_percentage_error(test['Log'], predictions2)`

Out[30]: 0.004140185361718522

In [31]: `mean_absolute_percentage_error(test['Thousands of Passengers'], np.exp(predictions2))`

Out[31]: 0.02596284258776781

In [ ]: